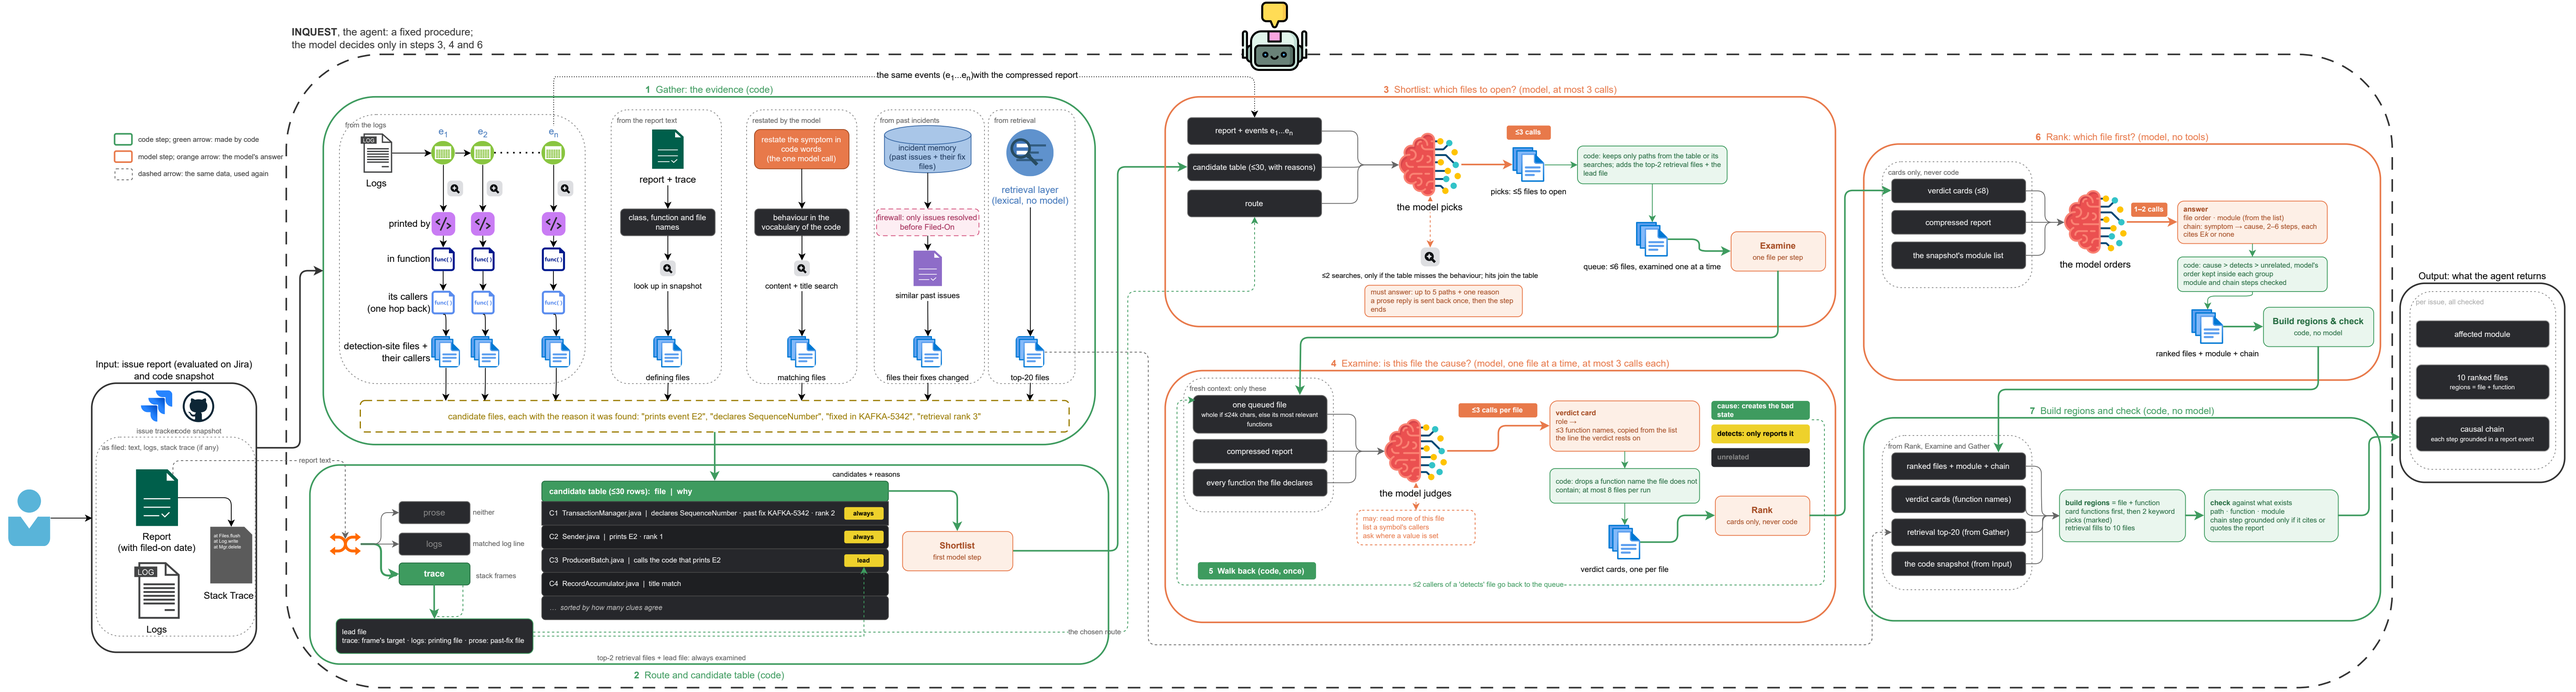

INQUEST, the agent: a fixed procedure;
the model decides only in steps 3, 4 and 6
code step; green arrow: made by code
model step; orange arrow: the model's answer
dashed arrow: the same data, used again
Input: issue report (evaluated on Jira)
and code snapshot
issue tracker
code snapshot
as filed: text, logs, stack trace (if any)
Report
(with filed-on date)
Stack Trace
Logs
report text
1 Gather: the evidence (code)
from the logs
Logs
e1
e2
en
printed by
in function
its callers
(one hop back)
detection-site files +
their callers
from the report text
report + trace
class, function and file names
look up in snapshot
defining files
restated by the model
restate the symptom in code words
(the one model call)
behaviour in the vocabulary of the code
content + title search
matching files
from past incidents
incident memory
(past issues + their fix files)
firewall: only issues resolved before Filed-On
similar past issues
files their fixes changed
from retrieval
retrieval layer
(lexical, no model)
top-20 files
candidate files, each with the reason it was found: "prints event E2", "declares SequenceNumber", "fixed in KAFKA-5342", "retrieval rank 3"
the same events (e1…en) with the compressed report
2 Route and candidate table (code)
candidates + reasons
candidate table (≤30 rows): file | why
C1 TransactionManager.java | declares SequenceNumber · past fix KAFKA-5342 · rank 2
always
C2 Sender.java | prints E2 · rank 1
always
C3 ProducerBatch.java | calls the code that prints E2
lead
C4 RecordAccumulator.java | title match
… sorted by how many clues agree
prose
neither
logs
matched log line
trace
stack frames
lead file
trace: frame's target · logs: printing file · prose: past-fix file
Shortlist
first model step
the chosen route
top-2 retrieval files + lead file: always examined
3 Shortlist: which files to open? (model, at most 3 calls)
report + events e1…en
candidate table (≤30, with reasons)
route
the model picks
≤3 calls
picks: ≤5 files to open
code: keeps only paths from the table or its
searches; adds the top-2 retrieval files + the
lead file
queue: ≤6 files, examined one at a time
Examine
one file per step
≤2 searches, only if the table misses the behaviour; hits join the table
must answer: up to 5 paths + one reason
a prose reply is sent back once, then the step
ends
4 Examine: is this file the cause? (model, one file at a time, at most 3 calls each)
fresh context: only these
one queued file
whole if ≤24k chars, else its most relevant
functions
compressed report
every function the file declares
the model judges
≤3 calls per file
verdict card
role →
≤3 function names, copied from the list
the line the verdict rests on
code: drops a function name the file does not
contain; at most 8 files per run
cause: creates the bad state
detects: only reports it
unrelated
Rank
cards only, never code
verdict cards, one per file
may: read more of this file
list a symbol's callers
ask where a value is set
5 Walk back (code, once)
≤2 callers of a 'detects' file go back to the queue
6 Rank: which file first? (model, no tools)
cards only, never code
verdict cards (≤8)
compressed report
the snapshot's module list
the model orders
1–2 calls
answer
file order · module (from the list)
chain: symptom → cause, 2–6 steps, each
cites Ex or none
code: cause > detects > unrelated, model's
order kept inside each group
module and chain steps checked
ranked files + module + chain
Build regions & check
code, no model
7 Build regions and check (code, no model)
from Rank, Examine and Gather
ranked files + module + chain
verdict cards (function names)
retrieval top-20 (from Gather)
the code snapshot (from Input)
build regions = file + function
card functions first, then 2 keyword
picks (marked)
retrieval fills to 10 files
check against what exists
path · function · module
chain step grounded only if it cites or
quotes the report
Output: what the agent returns
per issue, all checked
affected module
10 ranked files
regions + file + function
causal chain
each step grounded in a report event